

```
In [ ]: import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import cv2
import os
from PIL import Image
from mlxtend.preprocessing import TransactionEncoder
from mlxtend.frequent_patterns import apriori, association_rules
import seaborn as sns
import tensorflow.keras
from tensorflow.keras import layers, models, optimizers
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from sklearn.model_selection import train_test_split
```

WARNING:tensorflow:From c:\Users\kaasa\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

Question 1

```
In [ ]: df = pd.read_csv("../Assignment3/Grocery_Items_24.csv")
```

```
In [ ]: # referred from https://rasbt.github.io/mlxtend/user_guide/frequent_patterns/association_rules/
transactions = []
for i in range(len(df)):
    transactions.append([str(item) for item in df.iloc[i, :]])

te = TransactionEncoder()
te_ary = te.fit(transactions).transform(transactions)
df_encoded = pd.DataFrame(te_ary, columns=te.columns_)

min_support = 0.01
frequent_itemsets = apriori(df_encoded, min_support=min_support, use_colnames=True)
```

```
# Generating association rules
min_confidence = 0.1
rules = association_rules(frequent_itemsets, metric="confidence", min_threshold=min_confidence)

# Print the association rules
print("Association Rules:")
print(rules[['antecedents', 'consequents', 'support', 'confidence']])
```

Association Rules:

	antecedents	consequents	support	confidence
0	(UHT-milk)	(nan)	0.020875	1.000000
1	(beef)	(nan)	0.034250	1.000000
2	(berries)	(nan)	0.021250	1.000000
3	(beverages)	(nan)	0.016750	0.992593
4	(bottled beer)	(nan)	0.045250	1.000000
..
81	(rolls/buns, nan)	(whole milk)	0.013500	0.120805
82	(rolls/buns)	(whole milk, nan)	0.013500	0.120670
83	(whole milk, soda)	(nan)	0.011375	1.000000
84	(soda, nan)	(whole milk)	0.011375	0.121495
85	(soda)	(whole milk, nan)	0.011375	0.121495

[86 rows x 4 columns]

```
In [ ]: def count_rules(min_support, min_confidence):
    frequent_itemsets = apriori(df_encoded, min_support=min_support, use_colnames=True)
    rules = association_rules(frequent_itemsets, metric="confidence", min_threshold=min_confidence)

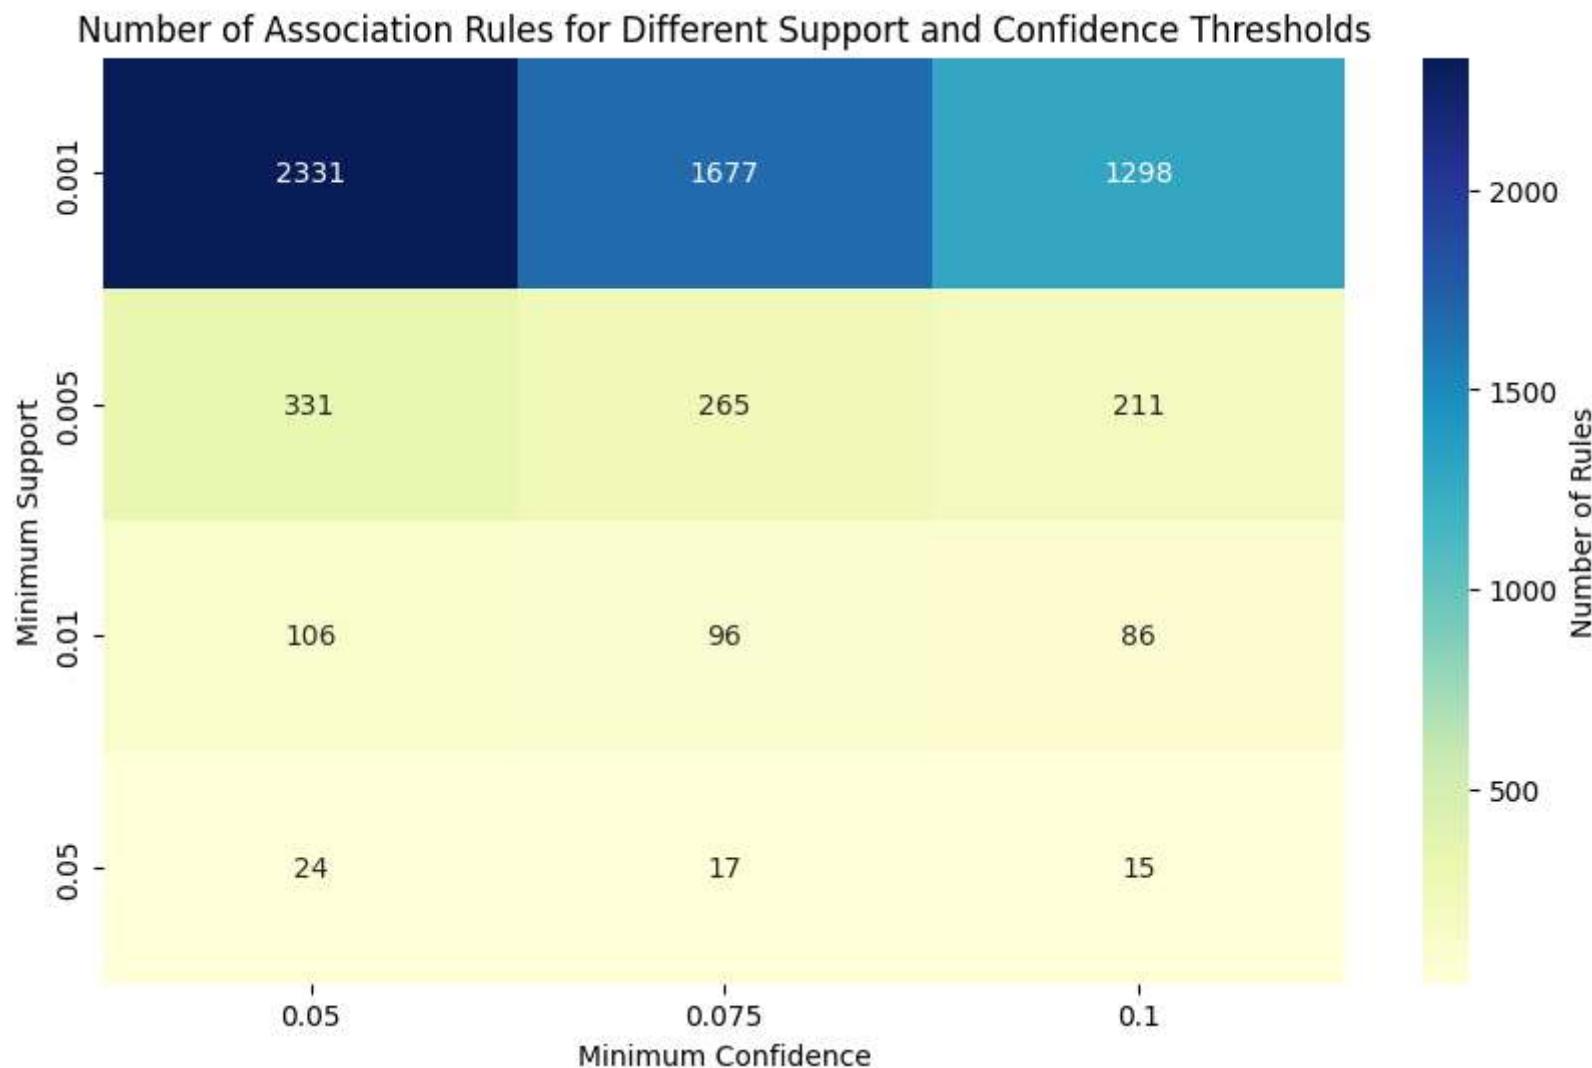
    return len(rules),rules

# the range of min_support values and min_confidence thresholds given
min_support_values = [0.001, 0.005, 0.01, 0.05]
min_confidence_thresholds = [0.05, 0.075, 0.1]
assoc_rules = pd.DataFrame(index=min_support_values, columns=min_confidence_thresholds)
results = pd.DataFrame(index=min_support_values, columns=min_confidence_thresholds)

# Iterate over min_support values and min_confidence thresholds to count the number of rules
for min_support in min_support_values:
    for min_confidence in min_confidence_thresholds:
        num_rules, rules = count_rules(min_support, min_confidence)
        results.at[min_support, min_confidence] = num_rules
        assoc_rules.at[min_support, min_confidence] = rules
```

```
# References:  
# - Pandas Documentation: https://pandas.pydata.org/pandas-docs/stable/  
# - mlxtend Documentation: http://rasbt.github.io/mlxtend/
```

```
In [ ]: # reffered from https://seaborn.pydata.org/generated/seaborn.heatmap.html  
results = results.apply(pd.to_numeric)  
  
plt.figure(figsize=(10, 6))  
sns.heatmap(results, annot=True, cmap="YlGnBu", fmt='g', cbar_kws={'label': 'Number of Rules'})  
plt.title('Number of Association Rules for Different Support and Confidence Thresholds')  
plt.xlabel('Minimum Confidence')  
plt.ylabel('Minimum Support')  
  
plt.show()
```



```
In [ ]: min_support = 0.005
frequent_itemsets = apriori(df_encoded, min_support=min_support, use_colnames=True)

rules = association_rules(frequent_itemsets, metric="confidence", min_threshold=0.0)

filtered_rules = rules[rules['support'] >= min_support]

max_confidence_rule = filtered_rules[filtered_rules['confidence'] == filtered_rules['confidence'].max()]
```

```

print("Association Rule with the Highest Confidence and Minimum Support 0.005:")
print(max_confidence_rule[['antecedents', 'consequents', 'support', 'confidence']])


# References:
# - Pandas Documentation: https://pandas.pydata.org/pandas-docs/stable/
# - mlxtend Documentation: http://rasbt.github.io/mlxtend/

```

Association Rule with the Highest Confidence and Minimum Support 0.005:

	antecedents	consequents	support	confidence
1	(UHT-milk)	(nan)	0.020875	1.0
2	(baking powder)	(nan)	0.008500	1.0
5	(beef)	(nan)	0.034250	1.0
7	(berries)	(nan)	0.021250	1.0
10	(bottled beer)	(nan)	0.045250	1.0
..
466	(whole milk, soda)	(nan)	0.011375	1.0
472	(yogurt, soda)	(nan)	0.006625	1.0
479	(whole milk, tropical fruit)	(nan)	0.008125	1.0
485	(yogurt, tropical fruit)	(nan)	0.006000	1.0
490	(whole milk, yogurt)	(nan)	0.009875	1.0

[113 rows x 4 columns]

The confidence value is 1

Question 2

```

In [ ]: import os
import tensorflow as tf
from tensorflow.keras import layers, models
import matplotlib.pyplot as plt
from tensorflow.keras.preprocessing.image import ImageDataGenerator

# Label map and data directory
labmap = {0: "n02089078-black-and-tan_coonhound",
          1: "n02091831-Saluki",
          2: "n02092002-Scottish_deerhound",
          3: "n02095314-wire-haired_fox_terrier"}

```

```
data_directory = '../DataSet/ProcessedDatasets/'\n\n# parameters\nbatch_size = 8\nimage_size = (331, 331)\nepochs = 20\n\n# an ImageDataGenerator for data augmentation\ndatagen = ImageDataGenerator(\n    rescale=1./255,\n    validation_split=0.2 # 80% training, 20% validation\n)\n\n# data generators\ntrain_generator = datagen.flow_from_directory(\n    data_directory,\n    target_size=image_size,\n    batch_size=batch_size,\n    class_mode='categorical',\n    classes=list(labmap.values()),\n    subset='training'\n)\n\nvalidation_generator = datagen.flow_from_directory(\n    data_directory,\n    target_size=image_size,\n    batch_size=batch_size,\n    class_mode='categorical',\n    classes=list(labmap.values()),\n    subset='validation'\n)\n\n# References:\n# - TensorFlow Documentation: https://www.tensorflow.org/\n# - Keras Documentation: https://keras.io/\n# - Matplotlib Documentation: https://matplotlib.org/\n# - ImageDataGenerator Documentation: https://www.tensorflow.org/api\_docs/python/tf/keras/preprocessing/image/ImageDataGenerator\n# - Python os Module Documentation: https://docs.python.org/3/library/os.html
```

Found 600 images belonging to 4 classes.
Found 148 images belonging to 4 classes.

3,3 Filter Size

```
In [ ]: model0 = models.Sequential()

# Convolutional Layer with 8 3 × 3 filters
model0.add(layers.Conv2D(8, (3, 3), activation='relu', input_shape=(331, 331, 3)))

# 1 max pooling with 2 × 2 pool size
model0.add(layers.MaxPooling2D((2, 2)))

# Flatten the Tensor
model0.add(layers.Flatten())

# 1 hidden Layer with 16 nodes for a fully connected neural network
model0.add(layers.Dense(16, activation='relu'))

# Output Layer has 4 nodes (since 4 classes) using 'softmax' activation function
model0.add(layers.Dense(4, activation='softmax'))

# Compile the model
model0.compile(optimizer='adam',
               loss='categorical_crossentropy',
               metrics=['accuracy'])

# Train the model
history = model0.fit(
    train_generator,
    steps_per_epoch=train_generator.samples // batch_size,
    epochs=epochs,
    validation_data=validation_generator,
    validation_steps=validation_generator.samples // batch_size
)

# Plot the training and validation accuracy
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```

```
# References:  
# - TensorFlow Documentation: https://www.tensorflow.org/  
# - Keras Documentation: https://keras.io/  
# - Matplotlib Documentation: https://matplotlib.org/
```

WARNING:tensorflow:From c:\Users\kaasa\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\backend.py:873: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From c:\Users\kaasa\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\layers\pooling\max_pooling2d.py:161: The name tf.nn.max_pool is deprecated. Please use tf.nn.max_pool2d instead.

WARNING:tensorflow:From c:\Users\kaasa\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\optimizers__init__.py:309: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

Epoch 1/20

WARNING:tensorflow:From c:\Users\kaasa\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\utils\tf_utils.py:492: The name tf.ragged.RaggedTensorValue is deprecated. Please use tf.compat.v1.ragged.RaggedTensorValue instead.

WARNING:tensorflow:From c:\Users\kaasa\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\engine\base_layer_utils.py:384: The name tf.executing_eagerly_outside_functions is deprecated. Please use tf.compat.v1.executing_eagerly_outside_functions instead.

75/75 [=====] - 3s 39ms/step - loss: 2.9650 - accuracy: 0.4017 - val_loss: 1.1885 - val_accuracy: 0.5069

Epoch 2/20

75/75 [=====] - 3s 37ms/step - loss: 1.1231 - accuracy: 0.4817 - val_loss: 1.1454 - val_accuracy: 0.4722

Epoch 3/20

75/75 [=====] - 3s 37ms/step - loss: 0.9529 - accuracy: 0.5917 - val_loss: 1.1178 - val_accuracy: 0.4931

Epoch 4/20

75/75 [=====] - 3s 37ms/step - loss: 0.7962 - accuracy: 0.6867 - val_loss: 1.1087 - val_accuracy: 0.5417

Epoch 5/20

75/75 [=====] - 3s 37ms/step - loss: 0.7365 - accuracy: 0.6867 - val_loss: 1.3236 - val_accuracy: 0.4375

Epoch 6/20

75/75 [=====] - 3s 37ms/step - loss: 0.6890 - accuracy: 0.7200 - val_loss: 1.1400 - val_accuracy: 0.4931

Epoch 7/20

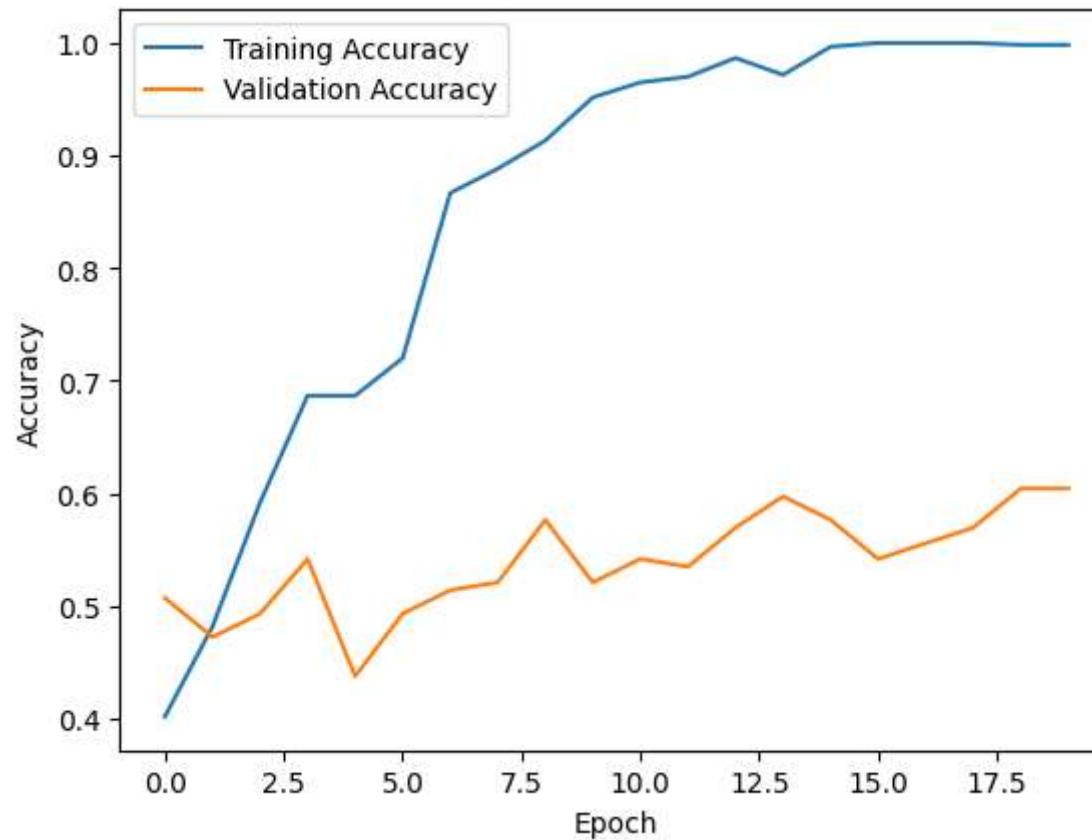
75/75 [=====] - 3s 37ms/step - loss: 0.4758 - accuracy: 0.8667 - val_loss: 1.1906 - val_accuracy: 0.5139

Epoch 8/20

75/75 [=====] - 3s 37ms/step - loss: 0.4023 - accuracy: 0.8883 - val_loss: 1.2752 - val_accuracy: 0.5208

Epoch 9/20

```
75/75 [=====] - 3s 37ms/step - loss: 0.3609 - accuracy: 0.9133 - val_loss: 1.1574 - val_accuracy: 0.5764
Epoch 10/20
75/75 [=====] - 3s 36ms/step - loss: 0.2729 - accuracy: 0.9517 - val_loss: 1.5291 - val_accuracy: 0.5208
Epoch 11/20
75/75 [=====] - 3s 37ms/step - loss: 0.2282 - accuracy: 0.9650 - val_loss: 1.4210 - val_accuracy: 0.5417
Epoch 12/20
75/75 [=====] - 3s 37ms/step - loss: 0.1907 - accuracy: 0.9700 - val_loss: 1.5635 - val_accuracy: 0.5347
Epoch 13/20
75/75 [=====] - 3s 36ms/step - loss: 0.1409 - accuracy: 0.9867 - val_loss: 1.5129 - val_accuracy: 0.5694
Epoch 14/20
75/75 [=====] - 3s 39ms/step - loss: 0.1542 - accuracy: 0.9717 - val_loss: 1.2778 - val_accuracy: 0.5972
Epoch 15/20
75/75 [=====] - 3s 37ms/step - loss: 0.1052 - accuracy: 0.9967 - val_loss: 1.5093 - val_accuracy: 0.5764
Epoch 16/20
75/75 [=====] - 3s 38ms/step - loss: 0.0697 - accuracy: 1.0000 - val_loss: 1.8391 - val_accuracy: 0.5417
Epoch 17/20
75/75 [=====] - 3s 39ms/step - loss: 0.0560 - accuracy: 1.0000 - val_loss: 1.7517 - val_accuracy: 0.5556
Epoch 18/20
75/75 [=====] - 3s 37ms/step - loss: 0.0470 - accuracy: 1.0000 - val_loss: 1.8102 - val_accuracy: 0.5694
Epoch 19/20
75/75 [=====] - 3s 37ms/step - loss: 0.0474 - accuracy: 0.9983 - val_loss: 1.8246 - val_accuracy: 0.6042
Epoch 20/20
75/75 [=====] - 3s 37ms/step - loss: 0.0466 - accuracy: 0.9983 - val_loss: 1.8272 - val_accuracy: 0.6042
```



Rowan Banner ID: 916459193

(a) 5*5 Filter Size

In []:

```
model1 = models.Sequential()

# Convolutional Layer with 8 5 × 5 filters
model1.add(layers.Conv2D(8, (5, 5), activation='relu', input_shape=(331, 331, 3)))

# 1 max pooling with 2 × 2 pool size
model1.add(layers.MaxPooling2D((2, 2)))

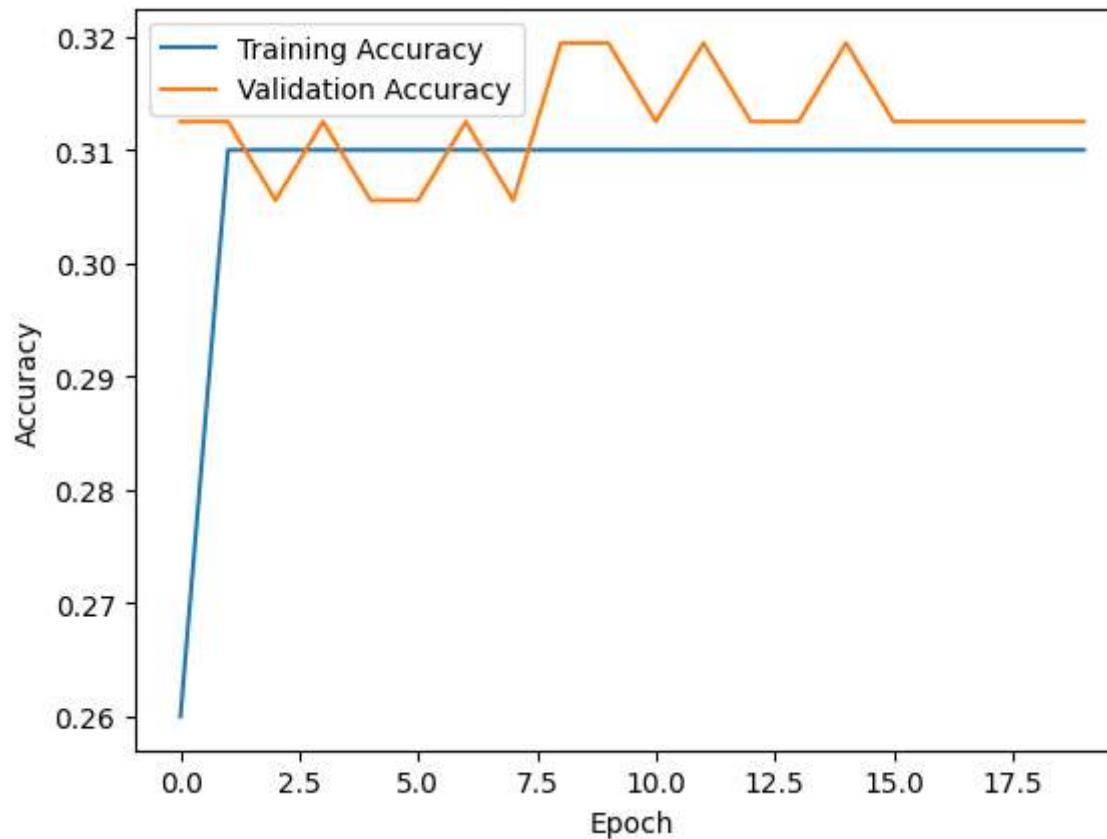
# Flatten the Tensor
model1.add(layers.Flatten())

# 1 hidden Layer with 16 nodes for a fully connected neural network
```

```
model1.add(layers.Dense(16, activation='relu'))  
  
# Output Layer has 4 nodes (since 4 classes) using 'softmax' activation function  
model1.add(layers.Dense(4, activation='softmax'))  
  
# Compile the model  
model1.compile(optimizer='adam',  
                loss='categorical_crossentropy',  
                metrics=['accuracy'])  
  
# Train the model  
history = model1.fit(  
    train_generator,  
    steps_per_epoch=train_generator.samples // batch_size,  
    epochs=epochs,  
    validation_data=validation_generator,  
    validation_steps=validation_generator.samples // batch_size  
)  
  
# Plot the training and validation accuracy  
plt.plot(history.history['accuracy'], label='Training Accuracy')  
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')  
plt.xlabel('Epoch')  
plt.ylabel('Accuracy')  
plt.legend()  
plt.show()  
  
# References:  
# - TensorFlow Documentation: https://www.tensorflow.org/  
# - Keras Documentation: https://keras.io/  
# - Matplotlib Documentation: https://matplotlib.org/
```

Epoch 1/20
75/75 [=====] - 3s 37ms/step - loss: 1.7335 - accuracy: 0.2600 - val_loss: 1.3847 - val_accuracy: 0.3125
Epoch 2/20
75/75 [=====] - 3s 36ms/step - loss: 1.3841 - accuracy: 0.3100 - val_loss: 1.3828 - val_accuracy: 0.3125
Epoch 3/20
75/75 [=====] - 3s 36ms/step - loss: 1.3820 - accuracy: 0.3100 - val_loss: 1.3808 - val_accuracy: 0.3056
Epoch 4/20
75/75 [=====] - 3s 35ms/step - loss: 1.3801 - accuracy: 0.3100 - val_loss: 1.3784 - val_accuracy: 0.3125
Epoch 5/20
75/75 [=====] - 3s 35ms/step - loss: 1.3786 - accuracy: 0.3100 - val_loss: 1.3776 - val_accuracy: 0.3056
Epoch 6/20
75/75 [=====] - 3s 35ms/step - loss: 1.3776 - accuracy: 0.3100 - val_loss: 1.3779 - val_accuracy: 0.3056
Epoch 7/20
75/75 [=====] - 3s 35ms/step - loss: 1.3765 - accuracy: 0.3100 - val_loss: 1.3741 - val_accuracy: 0.3125
Epoch 8/20
75/75 [=====] - 3s 35ms/step - loss: 1.3758 - accuracy: 0.3100 - val_loss: 1.3758 - val_accuracy: 0.3056
Epoch 9/20
75/75 [=====] - 3s 35ms/step - loss: 1.3753 - accuracy: 0.3100 - val_loss: 1.3712 - val_accuracy: 0.3194
Epoch 10/20
75/75 [=====] - 3s 36ms/step - loss: 1.3748 - accuracy: 0.3100 - val_loss: 1.3723 - val_accuracy: 0.3194
Epoch 11/20
75/75 [=====] - 3s 35ms/step - loss: 1.3744 - accuracy: 0.3100 - val_loss: 1.3737 - val_accuracy: 0.3125
Epoch 12/20
75/75 [=====] - 3s 36ms/step - loss: 1.3740 - accuracy: 0.3100 - val_loss: 1.3704 - val_accuracy: 0.3194
Epoch 13/20
75/75 [=====] - 3s 36ms/step - loss: 1.3739 - accuracy: 0.3100 - val_loss: 1.3721 - val_accuracy: 0.3125
Epoch 14/20
75/75 [=====] - 3s 36ms/step - loss: 1.3737 - accuracy: 0.3100 - val_loss: 1.3706 - val_accuracy: 0.3125

```
Epoch 15/20
75/75 [=====] - 3s 36ms/step - loss: 1.3735 - accuracy: 0.3100 - val_loss: 1.3696 - val_accuracy: 0.3194
Epoch 16/20
75/75 [=====] - 3s 36ms/step - loss: 1.3735 - accuracy: 0.3100 - val_loss: 1.3730 - val_accuracy: 0.3125
Epoch 17/20
75/75 [=====] - 3s 36ms/step - loss: 1.3733 - accuracy: 0.3100 - val_loss: 1.3728 - val_accuracy: 0.3125
Epoch 18/20
75/75 [=====] - 3s 36ms/step - loss: 1.3733 - accuracy: 0.3100 - val_loss: 1.3714 - val_accuracy: 0.3125
Epoch 19/20
75/75 [=====] - 3s 36ms/step - loss: 1.3732 - accuracy: 0.3100 - val_loss: 1.3714 - val_accuracy: 0.3125
Epoch 20/20
75/75 [=====] - 3s 36ms/step - loss: 1.3733 - accuracy: 0.3100 - val_loss: 1.3712 - val_accuracy: 0.3125
```



7*7 Filter Size

```
In [ ]: model2 = models.Sequential()

# Convolutional Layer with 8 7 × 7 filters
model2.add(layers.Conv2D(8, (7, 7), activation='relu', input_shape=(331, 331, 3)))

# 1 max pooling with 2 × 2 pool size
model2.add(layers.MaxPooling2D((2, 2)))

# Flatten the Tensor
model2.add(layers.Flatten())

# 1 hidden Layer with 16 nodes for a fully connected neural network
model2.add(layers.Dense(16, activation='relu'))
```

```
# Output Layer has 4 nodes (since 4 classes) using 'softmax' activation function
model2.add(layers.Dense(4, activation='softmax'))

# Compile the model
model2.compile(optimizer='adam',
                loss='categorical_crossentropy',
                metrics=['accuracy'])

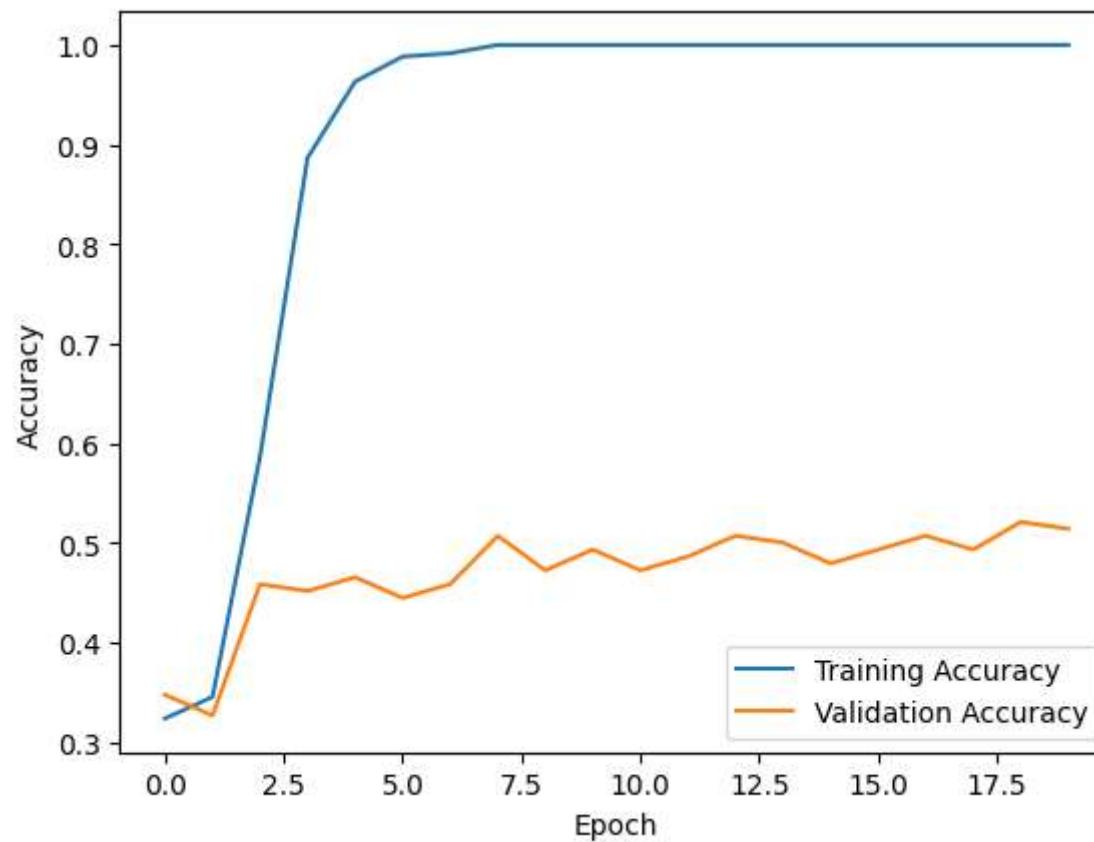
# Train the model
history = model2.fit(
    train_generator,
    steps_per_epoch=train_generator.samples // batch_size,
    epochs=epochs,
    validation_data=validation_generator,
    validation_steps=validation_generator.samples // batch_size
)

# Plot the training and validation accuracy
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

# References:
# - TensorFlow Documentation: https://www.tensorflow.org/
# - Keras Documentation: https://keras.io/
# - Matplotlib Documentation: https://matplotlib.org/
```

Epoch 1/20
75/75 [=====] - 3s 39ms/step - loss: 2.5722 - accuracy: 0.3233 - val_loss: 1.3670 - val_accuracy: 0.3472
Epoch 2/20
75/75 [=====] - 3s 37ms/step - loss: 1.3030 - accuracy: 0.3450 - val_loss: 1.3291 - val_accuracy: 0.3264
Epoch 3/20
75/75 [=====] - 3s 38ms/step - loss: 0.9599 - accuracy: 0.5867 - val_loss: 1.3747 - val_accuracy: 0.4583
Epoch 4/20
75/75 [=====] - 3s 38ms/step - loss: 0.4348 - accuracy: 0.8867 - val_loss: 1.4244 - val_accuracy: 0.4514
Epoch 5/20
75/75 [=====] - 3s 37ms/step - loss: 0.1729 - accuracy: 0.9633 - val_loss: 2.1354 - val_accuracy: 0.4653
Epoch 6/20
75/75 [=====] - 3s 37ms/step - loss: 0.0601 - accuracy: 0.9883 - val_loss: 2.1604 - val_accuracy: 0.4444
Epoch 7/20
75/75 [=====] - 3s 37ms/step - loss: 0.0324 - accuracy: 0.9917 - val_loss: 2.6269 - val_accuracy: 0.4583
Epoch 8/20
75/75 [=====] - 3s 37ms/step - loss: 0.0211 - accuracy: 1.0000 - val_loss: 2.6990 - val_accuracy: 0.5069
Epoch 9/20
75/75 [=====] - 3s 37ms/step - loss: 0.0075 - accuracy: 1.0000 - val_loss: 2.6248 - val_accuracy: 0.4722
Epoch 10/20
75/75 [=====] - 3s 37ms/step - loss: 0.0037 - accuracy: 1.0000 - val_loss: 2.5683 - val_accuracy: 0.4931
Epoch 11/20
75/75 [=====] - 3s 37ms/step - loss: 0.0026 - accuracy: 1.0000 - val_loss: 2.8576 - val_accuracy: 0.4722
Epoch 12/20
75/75 [=====] - 3s 37ms/step - loss: 0.0019 - accuracy: 1.0000 - val_loss: 2.7529 - val_accuracy: 0.4861
Epoch 13/20
75/75 [=====] - 3s 38ms/step - loss: 0.0015 - accuracy: 1.0000 - val_loss: 2.7974 - val_accuracy: 0.5069
Epoch 14/20
75/75 [=====] - 3s 38ms/step - loss: 0.0012 - accuracy: 1.0000 - val_loss: 3.0074 - val_accuracy: 0.5000

```
Epoch 15/20
75/75 [=====] - 3s 37ms/step - loss: 0.0010 - accuracy: 1.0000 - val_loss: 3.0727 - val_accuracy: 0.4792
Epoch 16/20
75/75 [=====] - 3s 37ms/step - loss: 8.6346e-04 - accuracy: 1.0000 - val_loss: 3.1099 - val_accuracy: 0.4931
Epoch 17/20
75/75 [=====] - 3s 37ms/step - loss: 7.3399e-04 - accuracy: 1.0000 - val_loss: 3.1448 - val_accuracy: 0.5069
Epoch 18/20
75/75 [=====] - 3s 37ms/step - loss: 6.1140e-04 - accuracy: 1.0000 - val_loss: 3.2050 - val_accuracy: 0.4931
Epoch 19/20
75/75 [=====] - 3s 38ms/step - loss: 5.0768e-04 - accuracy: 1.0000 - val_loss: 3.1439 - val_accuracy: 0.5208
Epoch 20/20
75/75 [=====] - 3s 37ms/step - loss: 4.2353e-04 - accuracy: 1.0000 - val_loss: 3.3316 - val_accuracy: 0.5139
```



Observation

The first model is learning slowly where the training and validation loss at initial epochs is very low, which gradually increased. The first model with a 3x3 filter is close to just right with the training and validation accuracies of (99.83,60.42) respectively

The first model with a 5x5 filter is underfit with the training and validation accuracies of (31.00,31.29) respectively

The first model with a 7x7 filter is overfit with the training and validation accuracies of (100,51.39) respectively

Model can be overfit when the training accuracy is very high and the validation accuracy is very low.

References:

- Pandas Documentation: <https://pandas.pydata.org/pandas-docs/stable/>
- mlxtend Documentation: <http://rasbt.github.io/mlxtend/>
- TensorFlow Documentation: <https://www.tensorflow.org/>
- Keras Documentation: <https://keras.io/>
- Matplotlib Documentation: <https://matplotlib.org/>
- ImageDataGenerator Documentation:
https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/image/ImageDataGenerator